

Assessment and Academic Integrity in the Age of Generative AI

AutoGenBook benchmark corpus

1 Ethical principles

The use of generative AI in teaching should rest on five principles: transparency, privacy, fairness, integrity and accountability. Transparency means that teachers and students state openly when and how generative tools were used. Privacy means that personal data of students are never entered into a tool without a legal basis and adequate protection. Fairness requires that no student is disadvantaged because they cannot afford a paid tool. Integrity means that assessed work reflects the student's own learning. Accountability means that a human remains responsible for every decision that affects a student.

These principles apply equally to teachers. A teacher who uses a model to draft feedback or to grade short answers remains accountable for the grade and must be able to explain it.

2 Data protection

Within the European Union the processing of personal data is governed by the General Data Protection Regulation (GDPR). Student names, identification numbers, grades, submitted work and even writing style can be personal data. Before student data are processed by an external AI service, the university must have a legal basis for the processing and a data processing agreement with the provider.

Three practical rules follow from the regulation. The first is data minimisation: only the data that are strictly necessary for the task should be processed. The second is anonymisation: names and identifiers should be removed from texts before they are pasted into a tool. The third is documentation: when a new tool will process student data at scale, a data protection impact assessment (DPIA) should be carried out before deployment. Enterprise licences and locally hosted models make it much easier to meet these obligations than free consumer accounts, whose terms often allow the provider to use conversations for further training.

3 The EU AI Act

The European regulation on artificial intelligence, known as the AI Act, was adopted in 2024 and introduces obligations that depend on the risk of an AI system. Systems with unacceptable risk are prohibited, high-risk systems must meet strict requirements on data quality, documentation and human oversight, and systems with limited risk are subject mainly to transparency obligations.

Education is explicitly named among the high-risk areas. AI systems used to decide on admission to education, to evaluate learning outcomes, or to monitor students during tests are classified as high-risk. A chat assistant used by a teacher to draft exercises is not high-risk in itself, but the same technology becomes high-risk when it is used to assign grades automatically.

4 Categories of assignments

Instead of a single rule for the whole university, many institutions ask teachers to classify each assignment into one of three categories and to state the category in the assignment brief. Table 1 summarises the scheme.

Category	Meaning	Typical example
Prohibited	Generative tools may not be used at all	Closed-book written exam
Permitted with declaration	Tools may be used if the student declares how	Seminar essay
Required	Using a tool is part of the task	Critique of an AI-generated text

Table 1: Three categories of assignments

For the category permitted with declaration, the declaration should name the tool, describe what it was used for, and include the key prompts. A short statement in the syllabus explains the categories to students at the start of the semester and prevents misunderstandings later.

5 Citing generative AI

When a student uses text or ideas produced by a generative tool, the use should be acknowledged like any other source. The major citation styles now provide templates for this purpose. In APA style, for example, the reference names the company that developed the model as the author, the name and version of the model, the date of use and the address of the service. The exact prompt should be reported in the text or in an appendix, because the same model can produce different answers to the same question.

6 Why AI detectors are not reliable

Software that claims to detect text written by generative AI produces both false positives and false negatives. Independent evaluations have shown that detectors frequently flag texts written by non-native speakers as generated, and that light paraphrasing is enough to hide generated text from them. A detector score is therefore not evidence of misconduct and must never be the sole basis for an accusation. Universities that have studied the problem recommend redesigning assessment instead of relying on detection.

7 Redesigning assessment

Assessment that is resistant to misuse of generative tools shifts attention from the final product to the process of learning. Oral examinations and short oral defences of written work allow the teacher to verify understanding directly. Process portfolios, in which students submit drafts, notes and reflections at several points during the semester, make it possible to follow how the work developed. Tasks connected to local data, personal experience or recent in-class discussion are difficult to complete with a general-purpose model alone.

A course redesign should start from the learning outcomes: for each outcome the teacher asks whether the current assessment still provides evidence that the student has achieved

it when generative tools are freely available. If not, the task is changed, moved into the classroom, or complemented by an oral component.

8 Practical checklist

Before the semester starts, a teacher can use the following checklist: decide the category of every assignment; add a statement on generative AI to the syllabus; choose tools that are covered by an institutional licence; prepare an example of a correct declaration; and plan at least one activity that develops the students' ability to evaluate generated content critically. During the semester, the teacher collects feedback from students on how the rules work in practice and adjusts them for the next run of the course.